

JMeter Analysis

September 20, 2022

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from IPython.display import Image
```

```
[2]: def load_dataset(results):
    df = pd.read_csv(results,
        ↳usecols=['timeStamp', 'elapsed', 'success', 'bytes', 'Latency',
        ↳'IdleTime', 'Connect'])
    df['totalElapsed'] = df.elapsed.cumsum()
    df['throughput'] = ((df.index+1)/(df.totalElapsed/(df.index+1))*60000)
    return df
```

```
[3]: def load_summary(summary):
    df = pd.read_csv(summary)
    return pd.read_csv(summary)
```

```
[4]: add500 = load_dataset('Add/Add_500T.csv')
addSum500 = load_summary('Add/Add_500T_Summary.csv')

add1000 = load_dataset('Add/Add_1000T.csv')
addSum1000 = load_summary('Add/Add_1000T_Summary.csv')
```

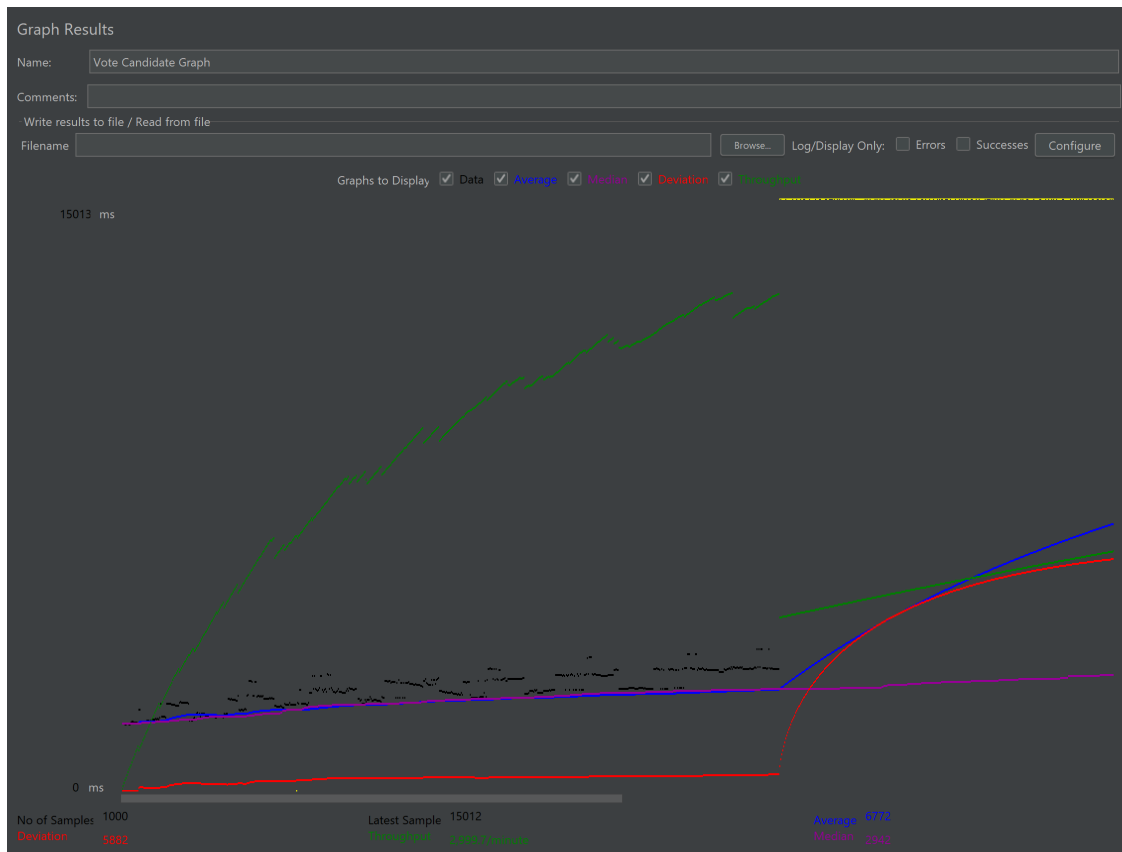
```
[5]: vote500 = load_dataset('Vote/Vote_500T.csv')
voteSum500 = load_summary('Vote/Vote_500T_Summary.csv')

vote1000 = load_dataset('Vote/Vote_1000T.csv')
voteSum1000 = load_summary('Vote/Vote_1000T_Summary.csv')
```

0.0.1 JMeter Graph Showing voteCandidate API Results: 1000 Threads

```
[6]: Image(filename='Vote/Vote_1000T.png')
```

```
[6]:
```



```
[7]: get500b = load_dataset('Get/Get_500T_50c.csv')
      getSum500b = load_summary('Get/Get_500T_50c_Summary.csv')

      get500_250c = load_dataset('Get/Get_500T_250c.csv')
      getSum500_250c = load_summary('Get/Get_500T_250c_Summary.csv')
```

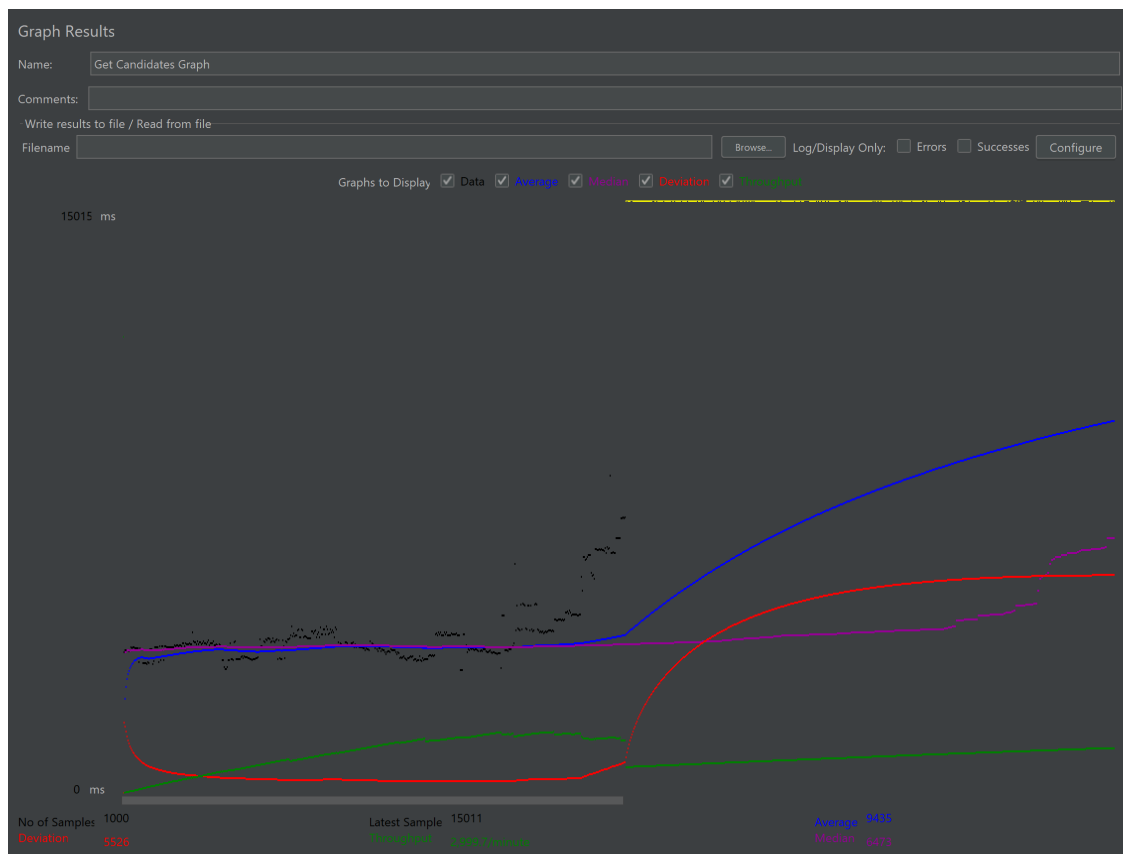
```
[8]: get1000b = load_dataset('Get/Get_1000T_50c.csv')
      getSum1000b = load_summary('Get/Get_1000T_50c_Summary.csv')

      get1000_250c = load_dataset('Get/Get_1000T_250c.csv')
      getSum1000_250c = load_summary('Get/Get_1000T_250c_Summary.csv')
```

0.0.2 JMeter Graph Showing getCandidates API Results: 1000 Threads 50 Candidates per pull

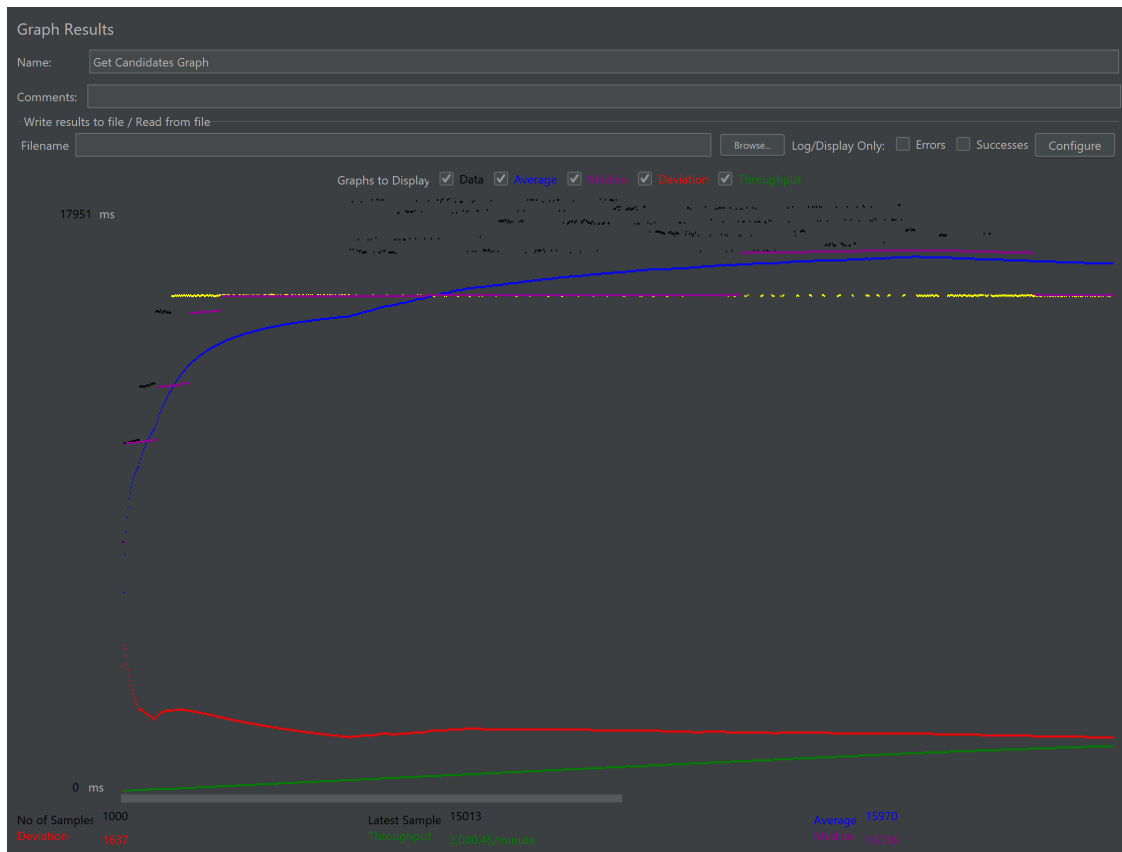
```
[9]: Image(filename='Get/Get_1000T_50c.png')
```

```
[9]:
```



[10]: `Image(filename='Get/Get_1000T_250c.png')`

[10]:



0.0.3 JMeter Graph Showing getCandidates API Results: 1000 Threads 250 Candidates per pull

[11]: `get1000b`

	timeStamp	elapsed	success	bytes	Latency	IdleTime	Connect	\
0	1663641500060	1753	False	2080	0	0	0	
1	1663641497543	3580	True	1703	3580	0	1	
2	1663641497549	3630	True	1703	3630	0	1	
3	1663641497548	3632	True	1703	3632	0	1	
4	1663641497555	3625	True	1703	3625	0	1	
..	
995	1663641502505	15015	False	2441	0	0	1	
996	1663641502516	15004	False	2441	0	0	0	
997	1663641502510	15010	False	2441	0	0	1	
998	1663641502521	15015	False	2441	0	0	0	
999	1663641502525	15011	False	2441	0	0	0	
	totalElapsed	throughput						
0	1753	34.227039						

1	5333	45.002813
2	8963	60.247685
3	12595	76.220723
4	16220	92.478422
..
995	9377415	6347.267344
996	9392419	6349.859392
997	9407429	6352.451876
998	9422444	6355.045464
999	9437455	6357.646209

[1000 rows x 9 columns]

[12]: get1000_250c

	timeStamp	elapsed	success	bytes	Latency	IdleTime	Connect	\
0	1663640344006	1753	False	2080	0	0	0	
1	1663640341471	7549	True	7654	7549	0	1	
2	1663640341476	10537	True	7654	10536	0	0	
3	1663640341476	10541	True	7654	10541	0	0	
4	1663640341481	10554	True	7654	10554	0	0	
..	
995	1663640346450	15002	False	2441	0	0	1	
996	1663640346445	15007	False	2441	0	0	1	
997	1663640346440	15012	False	2441	0	0	1	
998	1663640346436	15016	False	2441	0	0	1	
999	1663640346455	15013	False	2441	0	0	1	

	totalElapsed	throughput
0	1753	34.227039
1	9302	25.800903
2	19839	27.219114
3	30380	31.599737
4	40934	36.644354
..
995	15912570	3740.499492
996	15927577	3744.482918
997	15942589	3748.465196
998	15957605	3752.446561
999	15972618	3756.428658

[1000 rows x 9 columns]

[13]: add1000

	timeStamp	elapsed	success	bytes	Latency	IdleTime	Connect	\
0	1663634113661	1753	True	261	1753	0	1	

1	1663634113680	1735	True	261	1735	0	1
2	1663634113651	1766	True	261	1766	0	1
3	1663634113642	1775	True	261	1775	0	0
4	1663634113640	1777	True	261	1777	0	1
..
995	1663634118554	15000	False	2441	0	0	0
996	1663634118554	15000	False	2441	0	0	0
997	1663634118554	15000	False	2441	0	0	1
998	1663634118554	15000	False	2441	0	0	1
999	1663634118554	15000	False	2441	0	0	1

	totalElapsed	throughput
0	1753	34.227039
1	3488	68.807339
2	5254	102.778835
3	7029	136.577038
4	8806	170.338406
..
995	6775454	8784.792871
996	6790454	8782.997426
997	6805454	8781.227527
998	6820454	8779.483008
999	6835454	8777.763701

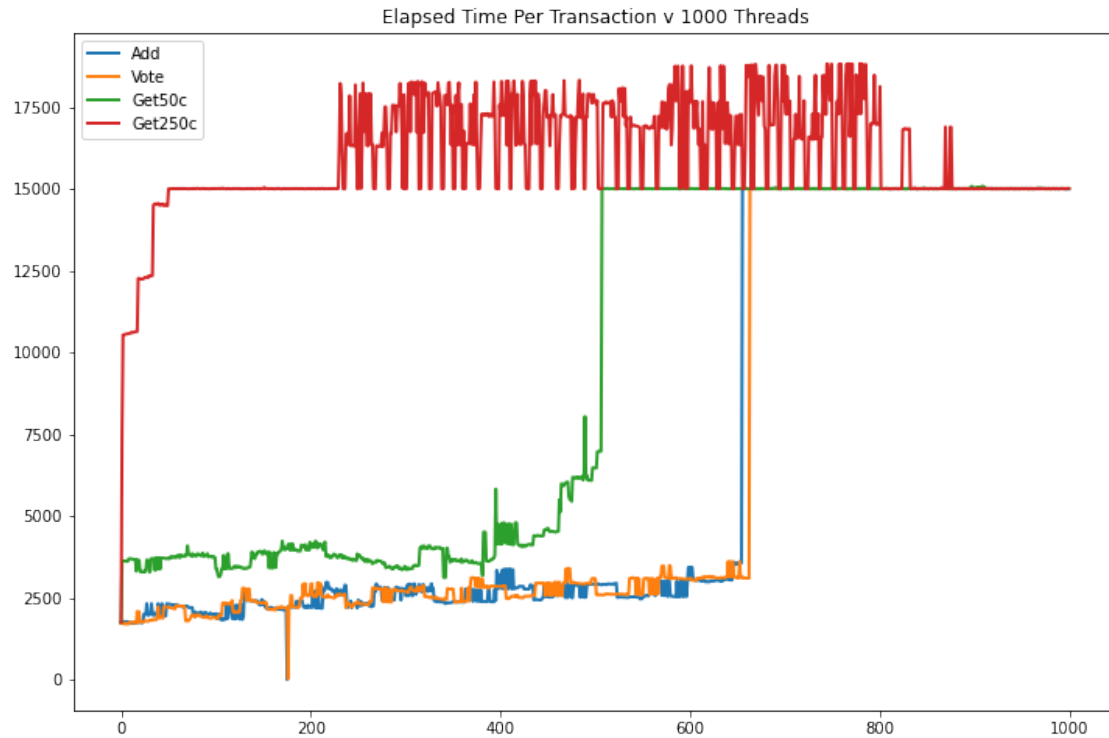
[1000 rows x 9 columns]

0.0.4 Elapsed Time per Transaction across 500 and 1000 Threads

In this section we plot the elapsed time of the add, vote, and get transactions across 500 and 1000 threads respectively. We can see that in the 500 thread plot, there is not a dramatic increase in elapsed time because the backend server does not hang. When we look at the 1000 thread plot we can see a near vertical line in the elapsed time, showing the number of threads that cause a hang in the backend server.


```
[14]: fig, ax = plt.subplots()
fig.set_figheight(8)
fig.set_figwidth(12)
ax.plot(range(0,1000), add1000.elapsed, linewidth=2.0, label="Add")
ax.plot(range(0,1000), vote1000.elapsed, linewidth=2.0, label="Vote")
ax.plot(range(0,1000), get1000b.elapsed, linewidth=2.0, label="Get50c")
ax.plot(range(0,1000), get1000_250c.elapsed, linewidth=2.0, label="Get250c")
ax.title.set_text('Elapsed Time Per Transaction v 1000 Threads')
ax.legend(loc="upper left")
```

```
[14]: <matplotlib.legend.Legend at 0x24fc9de68e0>
```



```
[15]: fig, ax = plt.subplots()
fig.set_figheight(8)
fig.set_figwidth(12)
ax.plot(range(0,500), add500.elapsed, linewidth=2.0, label="Add")
ax.plot(range(0,500), vote500.elapsed, linewidth=2.0, label="Vote")
ax.plot(range(0,500), get500b.elapsed, linewidth=2.0, label="Get 50c")
ax.plot(range(0,500), get500_250c.elapsed, linewidth=2.0, label="Get 250c")
ax.title.set_text('Elapsed Time Per Transaction v 500 Threads')
ax.legend(loc="upper left")
```

```
[15]: <matplotlib.legend.Legend at 0x24fc9e85340>
```




0.0.5 Throughput Per Minute over 1000 Threads

In the following section we calculate the average throughput per minute of each transaction across each the add, vote, and get function. We test the getfunction while retrieving 50 candidates as opposed to 250 candidates. Throughput is calculated as (number of requests)/(total elapsed time). Similar to the plot of elapsed time for 1000 threads we can see a dip in throughput for the add, vote, and get 50 candidates functions when the backend hangs. In comparison we do not see a throughput dip for retrieval of 250 candidates, since the throughput was much slower in growth.

```
[16]: fig, ax = plt.subplots()
fig.set_figheight(8)
fig.set_figwidth(12)
ax.plot(range(0,1000), add1000.throughput, linewidth=2.0, label="Add")
ax.plot(range(0,1000), vote1000.throughput, linewidth=2.0, label="Vote")
ax.plot(range(0,1000), get1000b.throughput, linewidth=2.0, label="Get 50c")
ax.plot(range(0,1000), get1000_250c.throughput, linewidth=2.0, label="Get 250c")
ax.title.set_text('Throughput/Min v Threads')
ax.legend(loc="upper left")
```

```
[16]: <matplotlib.legend.Legend at 0x24fca38bca0>
```

[17]: addSum1000

```
[17]:      Label  # Samples  Average  Min    Max  Std. Dev.  Error %  \
0  HTTP Request      1000    6835    3  15028    5941.17  34.600%
1      TOTAL          1000    6835    3  15028    5941.17  34.600%
```

	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
0	49.77601	49.33	7.34	1014.9
1	49.77601	49.33	7.34	1014.9

[18]: voteSum1000

```
[18]:      Label  # Samples  Average  Min    Max  Std. Dev.  Error %  \
0  HTTP Request      1000    6772    20  15019    5882.95  33.800%
1      TOTAL          1000    6772    20  15019    5882.95  33.800%
```

	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
0	49.995	48.7	7.27	997.5
1	49.995	48.7	7.27	997.5

[19]: getSum1000b

```
[19]:
```

	Label	# Samples	Average	Min	Max	Std. Dev.	Error %	\
0	HTTP Request	1000	9737	3	15028	5212.69	49.500%	
1	TOTAL	1000	9737	3	15028	5212.69	49.500%	

	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
0	49.9975	100.95	5.2	2067.6
1	49.9975	100.95	5.2	2067.6

```
[20]: getSum1000_250c
```

```
[20]:
```

	Label	# Samples	Average	Min	Max	Std. Dev.	Error %	\
0	HTTP Request	1000	15970	2	18835	1637.37	49.100%	
1	TOTAL	1000	15970	2	18835	1637.37	49.100%	

	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
0	50.0075	248.77	5.24	5094.1
1	50.0075	248.77	5.24	5094.1

0.0.6 Throughput per Second


In the following segment we plot throughput per second of each function in both the 500 thread and 1000 thread groups respectively. In this we actually see that when testing on our local machine that there is an upper limit of approximately 50 throughput per second. This also tells us that on the 500 thread group the upper limit on the throughput is actually bottlenecked by the number of threads being called instead of any problems related to function or server.

```
[21]: r1 = ['500t', addSum500.Throughput[1], voteSum500.Throughput[1], getSum500b.
    ↪Throughput[1], getSum500_250c.Throughput[1]]
r2 = ['1000t', addSum1000.Throughput[1], voteSum1000.Throughput[1], getSum1000b.
    ↪Throughput[1], getSum1000_250c.Throughput[1]]

df = pd.DataFrame([r1, r2], columns=['Threads', 'Add', 'Vote', 'Get 50c', 'Get_
    ↪250c'])

df.plot(x='Threads',
        kind='bar',
        stacked=False,
        title='Throughput/Sec Across Thread Groups',
        figsize=(12, 8),
        rot=0)
```

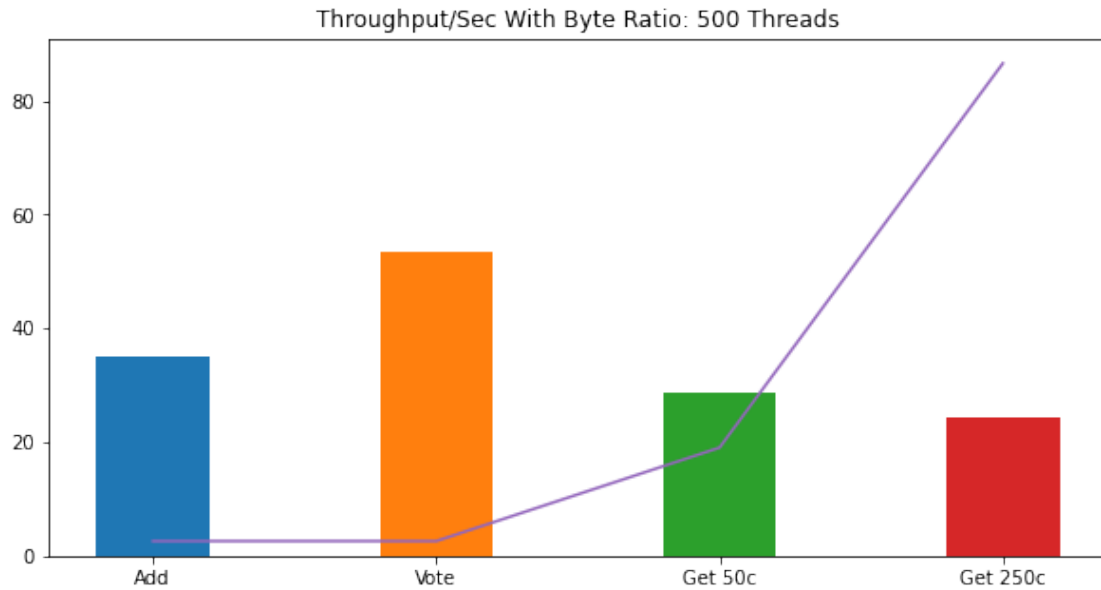
```
[21]: <AxesSubplot:title={'center': 'Throughput/Sec Across Thread Groups'},
      xlabel='Threads'>
```




0.0.7 Throughput vs Bytes

When comparing the results of the functions for 500 threads to 1000 threads, we can see that the number of bytes play a less significant role on throughput per second when there is a hang in the backend server. When the server is running optimally we can see that the increase in number of bytes can lead to lower throughput, however, when factoring in a server hang the number of bytes do not have a significant affect on overall throughput. This is also indicated by the plots in throughput per second where we see that the limiting factor on throughput is actually the number of threads being called.

```
[22]: Bytes = [addSum500['Avg. Bytes'][1], voteSum500['Avg. Bytes'][1],
↳ getSum500b['Avg. Bytes'][1], getSum500_250c['Avg. Bytes'][1]]
Throughput = [addSum500.Throughput[1], voteSum500.Throughput[1], getSum500b.
↳ Throughput[1], getSum500_250c.Throughput[1]]
labels = ['Add', 'Vote', 'Get 50c', 'Get 250c']
colors= ['tab:blue', 'tab:orange', 'tab:green', 'tab:red']
ScaledBytes=[x/100 for x in Bytes]
fig = plt.figure(figsize = (10, 5))
plt.bar(labels, Throughput, width=0.4, color = colors)
plt.plot(labels, ScaledBytes, color = 'tab:purple')
plt.title("Throughput/Sec With Byte Ratio: 500 Threads")
plt.show()
```



```
[23]: Bytes = [addSum1000['Avg. Bytes'][1], voteSum1000['Avg. Bytes'][1],
↳getSum1000b['Avg. Bytes'][1], getSum1000_250c['Avg. Bytes'][1]]
Throughput = [addSum1000.Throughput[1], voteSum1000.Throughput[1], getSum1000b.
↳Throughput[1], getSum1000_250c.Throughput[1]]
labels = ['Add', 'Vote', 'Get 50c', 'Get 250c']
colors= ['tab:blue', 'tab:orange', 'tab:green', 'tab:red']
ScaledBytes=[x/100 for x in Bytes]
fig = plt.figure(figsize = (10, 5))
plt.bar(labels, Throughput, width=0.4, color = colors)
plt.plot(labels, ScaledBytes, color = 'tab:purple')
plt.title("Throughput/Sec With Byte Ratio: 1000 Threads")
plt.show()
```



0.0.8 Comparison of Failed Transactions Per Function with Byte Ratio

After seeing that the average throughput for each of the functions is about the same when running the 1000 thread group we need to consider the number of transactions that failed. In this case we can see that the increase in bytes may lead to earlier cases where we cannot connect to the server. For this segment we can also refer to the previous graph of elapsed time per transaction, increasing the number of return bytes per transaction increases the elapsed time per transaction, which eventually leads to a bottleneck in network transaction time.

```
[30]: def calcFailed(df, summary):
        failed = df.success.value_counts()[0]
        summary['Failed'] = failed
        return summary
```

```
[34]: addSum1000 = calcFailed(add1000, addSum1000)
        voteSum1000 = calcFailed(vote1000, voteSum1000)
        getSum1000b = calcFailed(get1000b, getSum1000b)
        getSum1000_250c = calcFailed(get1000_250c, getSum1000_250c)
```

```
[42]: FailedTrans = [addSum1000['Failed'][1],
                    voteSum1000['Failed'][1],
                    getSum1000b['Failed'][1],
                    getSum1000_250c['Failed'][1]]
        Bytes = [addSum1000['Avg. Bytes'][1], voteSum1000['Avg. Bytes'][1],
        ↪ getSum1000b['Avg. Bytes'][1], getSum1000_250c['Avg. Bytes'][1]]
        ScaledBytes=[x/10 for x in Bytes]
```

```

labels = ['Add', 'Vote', 'Get 50c', 'Get 250c']
colors= ['tab:blue', 'tab:orange', 'tab:green', 'tab:red']

fig = plt.figure(figsize = (10, 5))
plt.bar(labels, FailedTrans, width=0.4, color = colors)
plt.plot(labels, ScaledBytes, color = 'tab:purple')

plt.title("Failed Transactions With Byte Ratio: 1000 Threads")
plt.show()

```

